In [2]:
```
# model optimization
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
```

In [3]:
```
# reading the dataset
data = pd.read_csv("C:\\Users\\hm\\Desktop\\LINEAR PROGRAMMING CENTRALS\\studer
```

In [4]:
```
data
```

Out[4]:

	Study Hours	Exam Scores
0	3.7	87.9
1	9.5	143.6
2	7.3	123.7
3	6.0	99.9
4	1.6	64.5
...	...	...
95	4.9	95.3
96	5.2	101.9
97	4.3	94.5
98	0.3	53.9
99	1.1	64.9

100 rows × 2 columns

In [5]:
```
x = np.array(data["Study Hours"]).reshape(-1,1)
y = np.array(data["Exam Scores"])
```

In [6]:
```
x
```

Out[6]:
```
array([[3.7],
       [9.5],
       [7.3],
       [6. ],
       [1.6],
       [1.6],
       [0.6],
       [8.7],
       [6. ],
       [7.1],
       [0.2],
       [9.7],
       [8.3],
       [2.1],
       [1.8],
       [1.8],
       [3. ],
       [5.2],
       [4.3],
       [2.9],
       [6.1],
       [1.4],
       [2.9],
       [3.7],
       [4.6],
       [7.9],
       [2. ],
       [5.1],
       [5.9],
       [0.5],
       [6.1],
       [1.7],
       [0.7],
       [9.5],
       [9.7],
       [8.1],
       [3. ],
       [1. ],
       [6.8],
       [4.4],
       [1.2],
       [5. ],
       [0.3],
       [9.1],
       [2.6],
       [6.6],
       [3.1],
       [5.2],
       [5.5],
       [1.8],
       [9.7],
       [7.8],
       [9.4],
       [8.9],
       [6. ],
       [9.2],
       [0.9],
```

```
[2. ],
[0.5],
[3.3],
[3.9],
[2.7],
[8.3],
[3.6],
[2.8],
[5.4],
[1.4],
[8. ],
[0.7],
[9.9],
[7.7],
[2. ],
[0.1],
[8.2],
[7.1],
[7.3],
[7.7],
[0.7],
[3.6],
[1.2],
[8.6],
[6.2],
[3.3],
[0.6],
[3.1],
[3.3],
[7.3],
[6.4],
[8.9],
[4.7],
[1.2],
[7.1],
[7.6],
[5.6],
[7.7],
[4.9],
[5.2],
[4.3],
[0.3],
[1.1]])
```

In [7]:
```
y
```

Out[7]:
```
array([ 87.9, 143.6, 123.7,  99.9,  64.5,  67.4,  63.2, 134. , 106.1,
       118.3,  56.6, 148.6, 130.6,  73.8,  68.7,  73.2,  76.9, 100.8,
        91.2,  71.8, 112.7,  65.3,  79.2,  85.5,  88.5, 126.4,  68.3,
        97.4, 108.4,  56.7, 120.2,  67.9,  57.8, 144.5, 137. , 130.7,
        80.8,  72.1, 117.5,  95.5,  62. ,  93.7,  59.2, 144.7,  79.8,
       111.7,  88.2,  95. , 107.6,  79.4, 142. , 124.7, 144.4, 137. ,
       102. , 142.5,  53.5,  72. ,  49.9,  90.3,  85. ,  75.5, 136.9,
        79.5,  79.2, 110.8,  56.1, 131.1,  58.8, 152.6, 121. ,  63.3,
        53.2, 133. , 121.9, 124.6, 123.7,  58.6,  87.3,  58. , 145.6,
       114.7,  77.1,  59.6,  76.2,  86.5, 128.8, 109.7, 143.5,  99.3,
        66.1, 130.8, 124.9, 102.4, 122.6,  95.3, 101.9,  94.5,  53.9,
        64.9])
```

In [8]:
```
data.isnull().sum()
```

Out[8]:
```
Study Hours    0
Exam Scores    0
dtype: int64
```

In [9]:
```
# visualizing the relationship between  study hours and exam score
import matplotlib.pyplot as plt
from matplotlib import style
```

In [10]:
```
# ploting the scatter graph
style.use("seaborn")
plt.scatter(x, y, label= "datapoints", color = "red", s = 80)
plt.title("Relationship between STUDY HOURS AND EXAM SCORE")
plt.xlabel("Study hours")
plt.ylabel("exam score")
```

```
C:\Users\hm\AppData\Local\Temp\ipykernel_9728\1976094932.py:2: MatplotlibDepr
ecationWarning: The seaborn styles shipped by Matplotlib are deprecated since
3.6, as they no longer correspond to the styles shipped by seaborn. However,
they will remain available as 'seaborn-v0_8-<style>'. Alternatively, directly
use the seaborn API instead.
  style.use("seaborn")
```

Out[10]:
```
Text(0, 0.5, 'exam score')
```

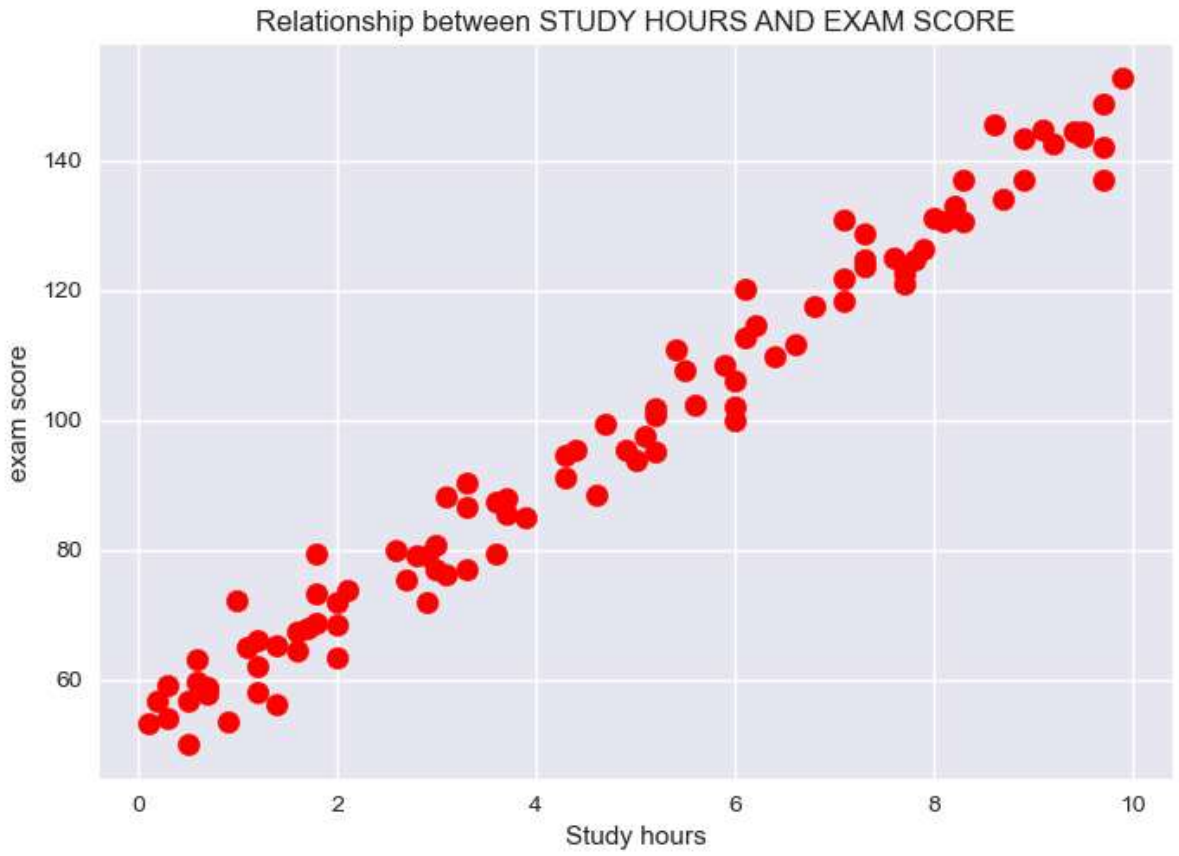

In [11]:
```
# spliting the dataset into train and test
x_train, x_test, y_train,y_test =  train_test_split(x,y, test_size = 0.2, rand
```

In [12]:
```
# data preprocessing
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

In [13]:
```
# building the linear regreesion model
model =LinearRegression()
model.fit(x_train,y_train)
```

Out[13]:
```
LinearRegression
LinearRegression()
```

In [14]:
```
y_pred = model.predict(x_test_scaled)
```

In [15]:
```
y_pred
```

Out[15]:
```
array([36.94787956, 64.87197905, 60.83475984, 57.13397558, 43.67657823,
       49.73240704, 44.68588303, 63.86267425, 35.60213983, 47.3773625 ,
       49.3959721 , 55.45180091, 62.51693451, 66.89058865, 38.96648916,
       40.3122289 , 60.83475984, 37.2843145 , 62.85336944, 40.64866383])
```

EVALUATING THE MODEL

In [16]:
```
MAE =mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)
R2 =r2_score(y_test, y_pred)
```

In [17]:
```
MAE
```

Out[17]:
```
47.715961961709375
```

In [18]:
```
MSE
```

Out[18]:
```
2682.9281679636674
```

In [19]:
```
R2
```

Out[19]:
```
-1.8659723784262328
```

In [20]:
```
model.intercept_
```

Out[20]:
```
50.691246598122966
```

In [21]:
```
model.coef_
```

Out[21]:
```
array([9.79642549])
```

In [22]:
```
# improving the model
y_pred_improved= model.predict(x_test_scaled)
MAE_improved = mean_absolute_error(y_test,y_pred_improved)
MSE_improved = mean_squared_error(y_test,y_pred_improved)
R2_improved = r2_score(y_test, y_pred_improved)
```

In [23]:
```
MAE_improved
```

Out[23]: 47.715961961709375

In [24]:
```
MSE_improved
```

Out[24]: 2682.9281679636674

In [ ]:
```
```